

Assignment 3 - Flight Delay Problem

Steve Spence

10/17/2019

Load Data Set and Libraries

First, we will load all of the packages that will be required for this problem. Specifically, “ISLR”, “caret”, “dplyr”, “e1071”, and “pROC” will be loaded for this problem.

Next, we will import the “FlightDelays” data set into the RStudio environment.

```
# Import data set from BlackBoard into the RStudio environment
```

```
FlightDelays <- read.csv("FlightDelays.csv")
```

Data Structure

A summary of the data set will be displayed to review the data set.

```
# Investigate the structure of the data set
```

```
str(FlightDelays)
```

```
## 'data.frame': 2201 obs. of 13 variables:
## $ CRS_DEP_TIME : int 1455 1640 1245 1715 1039 840 1240 1645 1715 2120
## $ CARRIER : Factor w/ 8 levels "CO","DH","DL",...: 5 2 2 2 2 2 2 2 2 2
## $ DEP_TIME : int 1455 1640 1245 1709 1035 839 1243 1644 1710 2129
## $ DEST : Factor w/ 3 levels "EWR","JFK","LGA": 2 2 3 3 3 2 2 2 2 2
## $ DISTANCE : int 184 213 229 229 229 228 228 228 228 228 ...
## $ FL_DATE : Factor w/ 31 levels "01/01/2004","01/02/2004",...: 1 1 1
## $ FL_NUM : int 5935 6155 7208 7215 7792 7800 7806 7810 7812 7814
## $ ORIGIN : Factor w/ 3 levels "BWI","DCA","IAD": 1 2 3 3 3 3 3 3 3 3
## $ Weather : int 0 0 0 0 0 0 0 0 0 0 ...
## $ DAY_WEEK : int 4 4 4 4 4 4 4 4 4 4 ...
## $ DAY_OF_MONTH : int 1 1 1 1 1 1 1 1 1 1 ...
## $ TAIL_NUM : Factor w/ 549 levels "N10323","N10575",...: 526 263 382
## $ Flight.Status: Factor w/ 2 levels "delayed","ontime": 2 2 2 2 2 2 2 2 2 2
```

```
# Investigate the summary of the data set
```

```
summary(FlightDelays)
```

```
##   CRS_DEP_TIME    CARRIER    DEP_TIME    DEST    DISTANCE
##   Min.   : 600    DH       :551    Min.   : 10    EWR: 665    Min.   :169.0
##   1st Qu.:1000    RU       :408    1st Qu.:1004    JFK: 386    1st Qu.:213.0
##   Median :1455    US       :404    Median :1450    LGA:1150    Median :214.0
##   Mean   :1372    DL       :388    Mean   :1369                Mean   :211.9
##   3rd Qu.:1710    MQ       :295    3rd Qu.:1709                3rd Qu.:214.0
##   Max.   :2130    CO       : 94    Max.   :2330                Max.   :229.0
##               (Other): 61
##   FL_DATE    FL_NUM    ORIGIN    Weather
##   1/22/2004 : 86    Min.   : 746    BWI: 145    Min.   :0.00000
##   01/06/2004: 85    1st Qu.:2156    DCA:1370    1st Qu.:0.00000
##   01/08/2004: 85    Median :2385    IAD: 686    Median :0.00000
##   1/13/2004 : 85    Mean   :3815                Mean   :0.01454
##   1/20/2004 : 85    3rd Qu.:6155                3rd Qu.:0.00000
##   1/21/2004 : 85    Max.   :7924                Max.   :1.00000
##   (Other)    :1690
##   DAY_WEEK    DAY_OF_MONTH    TAIL_NUM    Flight.Status
##   Min.   :1.000    Min.   : 1.00    N225DL : 65    delayed: 428
##   1st Qu.:2.000    1st Qu.: 8.00    N242DL : 56    ontime :1773
##   Median :4.000    Median :16.00    N223DZ : 50
##   Mean   :3.905    Mean   :16.02    N221DL : 45
##   3rd Qu.:5.000    3rd Qu.:23.00    N241DL : 36
##   Max.   :7.000    Max.   :31.00    N722UW : 36
##               (Other):1913
```

We can see from the data set above, it appears that 428 out of 2,201 flights (Approximately 19.5%) are delayed historically. This can be compared against the final model at the end for a reality check.

Given in the problem statement, we will only be using five predictors: “DAY_WEEK”, “DEP_TIME”, “ORIGIN”, “DEST”, and “CARRIER” along with the dependent variable “Flight.Status”. Therefore, we will re-write the dataframe with these 6 variables.

Additionally, we will need to convert all variables to factors for Naive Bayes model.

```
# Isolate the 6 variables previously discussed
```

```
FlightDelays <- FlightDelays[ , c(10, 1, 8, 4, 2, 13)]
```

```
# Review the new structure of the dataset
```

```
str(FlightDelays)
```

```
## 'data.frame':    2201 obs. of  6 variables:
##  $ DAY_WEEK      : int   4 4 4 4 4 4 4 4 4 4 ...
##  $ CRS_DEP_TIME  : int  1455 1640 1245 1715 1039 840 1240 1645 1715 2120
```

```
...
## $ ORIGIN      : Factor w/ 3 levels "BWI","DCA","IAD": 1 2 3 3 3 3 3 3 3
3 ...
## $ DEST        : Factor w/ 3 levels "EWR","JFK","LGA": 2 2 3 3 3 2 2 2 2
2 ...
## $ CARRIER    : Factor w/ 8 levels "CO","DH","DL",...: 5 2 2 2 2 2 2 2 2
2 ...
## $ Flight.Status: Factor w/ 2 levels "delayed","ontime": 2 2 2 2 2 2 2 2 2
2 ...
```

Now that we have the 6 variables in question, we must ensure all of them are converted to factors for proper use in the Naive Bayes algorithm. As shown above, “DAY_WEEK” and “CRS_DEP_TIME” are the two remaining variables that need to be converted to factors.

Additionally, “CRS_DEP_TIME” will also need to be converted to a factor with 16 time ranges (as stated in the initial problem statement). These time ranges will stretch from “6:00am to 7:00am” departure time to the final time range of “9:00pm to 10:00pm”. These will correlate to factor level 1 to 16, respectively.

```
# Convert "DAY_WEEK" to a factor for Naives Bayes
```

```
FlightDelays$DAY_WEEK <- as.factor(FlightDelays$DAY_WEEK)
```

```
# Convert "CRS_DEP_TIME" to a factor with 16 time ranges and Labels
```

```
FlightDelays$CRS_DEP_TIME <- cut(FlightDelays$CRS_DEP_TIME, breaks = c(600,
700, 800, 900, 1000, 1100, 1200, 1300, 1400, 1500, 1600, 1700, 1800, 1900,
2000, 2100, 2200))
```

```
# Verify the new structure of the data set
```

```
str(FlightDelays)
```

```
## 'data.frame':    2201 obs. of  6 variables:
## $ DAY_WEEK      : Factor w/ 7 levels "1","2","3","4",...: 4 4 4 4 4 4 4 4 4
4 ...
## $ CRS_DEP_TIME  : Factor w/ 16 levels "(600,700]","(700,800]","...: 9 11 7
12 5 3 7 11 12 16 ...
## $ ORIGIN       : Factor w/ 3 levels "BWI","DCA","IAD": 1 2 3 3 3 3 3 3 3
3 ...
## $ DEST         : Factor w/ 3 levels "EWR","JFK","LGA": 2 2 3 3 3 2 2 2 2
2 ...
## $ CARRIER     : Factor w/ 8 levels "CO","DH","DL",...: 5 2 2 2 2 2 2 2 2
2 ...
## $ Flight.Status: Factor w/ 2 levels "delayed","ontime": 2 2 2 2 2 2 2 2 2
2 ...
```

Data Preprocessing

To begin preprocessing data, we must split the data set into 60% training and 40% validation, per the problem description.

```
# Set the seed for randomized functions

set.seed(102019)

# Split the data into 60% training data and 40% validation data

FlightDelaysIndex <- createDataPartition(FlightDelays$DAY_WEEK, p=0.4, list =
F)

FlightDelaysValidation <- FlightDelays[FlightDelaysIndex,]

FlightDelaysTrain <- FlightDelays[-FlightDelaysIndex,]
```

Counts and Proportion Table by Airport

Assignment calls for a table containing a count and proportion of delayed flights by airport. For this, the “dplyr” package and summarise function will be used to return these values.

```
# Summarise flight delay statistics by airport

FlightDelays %>%
  group_by(Origin_Airport = ORIGIN, Flight_Status = Flight.Status) %>%
  summarise(Count_of_Flights = n()) %>%
  mutate(Proportion_for_Airport = 100*(Count_of_Flights /
sum(Count_of_Flights)))

## # A tibble: 6 x 4
## # Groups:   Origin_Airport [3]
##   Origin_Airport Flight_Status Count_of_Flights Proportion_for_Airport
##   <fct>          <fct>          <int>          <dbl>
## 1 BWI           delayed           37           25.5
## 2 BWI           ontime           108          74.5
## 3 DCA           delayed           221          16.1
## 4 DCA           ontime          1149          83.9
## 5 IAD           delayed           170          24.8
## 6 IAD           ontime           516          75.2
```

From this data table, we can see the total number of flights from each airport and the proportion of them that were delayed for each individual airport.

“BWI” had the lowest number of flights during this time period, but it also had the highest percentage of them being delayed.

Create Naive Bayes Model

Now that we have the training and validation data properly prepared, the Naive Bayes model can be created from the training data and then ran on the validation data.

```
# Create Naive Bayes model from the training data set
```

```
NB_Model <- naiveBayes(Flight.Status ~ ., data = FlightDelaysTrain)
```

```
# Use Model on the validation data set to predict if flights will be delayed
```

```
FlightDelaysValidation_Predicted <- predict(NB_Model, FlightDelaysValidation)
```

Confusion Matrix and ROC for Validation Data

A confusion matrix will be created to determine the accuracy of the model with varying statistics.

```
# Create confusion matrix for the Label outputs from the Naive Bayes model
```

```
CrossTable(x = FlightDelaysValidation$Flight.Status, y =  
FlightDelaysValidation_Predicted, prop.chisq = FALSE)
```

```
##
```

```
##
```

```
##      Cell Contents
```

```
## |-----|  
## |                      N |  
## |      N / Row Total |  
## |      N / Col Total |  
## |      N / Table Total |  
## |-----|
```

```
##
```

```
##
```

```
## Total Observations in Table:  883
```

```
##
```

```
##
```

```
##      FlightDelaysValidation$Flight.Status | FlightDelaysValidation_Predicted  
##      delayed      ontime      Row Total |  
## -----|-----|-----|-----|  
##      delayed      17      173      190 |  
##      0.089      0.911      0.215 |  
##      0.500      0.204 |  
##      0.019      0.196 |  
## -----|-----|-----|  
##      ontime      17      676      693 |  
##      0.025      0.975      0.785 |  
##      0.500      0.796 |  
##      0.019      0.766 |  
## -----|-----|-----|
```

##	Column Total	34	849	883
##		0.039	0.961	
##	-----	-----	-----	-----
##				
##				

From this confusion matrix, we can state that:

The sensitivity (true positive rate) of the model is: $676/693 = 97.5\%$ The specificity (true negative rate) of the model is: $17/190 = 9.0\%$

Next, there will be the creation of the ROC curve and return the AUC value for the validation data on this model.

Before the curve is plotted out, there must be another model ran with the raw probabilities listed.

```
# Re-run the model to return the raw probabilities
```

```
FlightDelaysValidation_Predicted2 <- predict(NB_Model,
FlightDelaysValidation, type = "raw")
```

```
# Return the first few values in the output table
```

```
head(FlightDelaysValidation_Predicted2)
```

```
##      delayed   ontime
## [1,] 0.13747017 0.8625298
## [2,] 0.20216762 0.7978324
## [3,] 0.36800677 0.6319932
## [4,] 0.29987116 0.7001288
## [5,] 0.03522145 0.9647786
## [6,] 0.09009041 0.9099096
```

Next, the curves can now be plotted out.

```
# Creating the ROC curve for the model
```

```
roc(FlightDelaysValidation$Flight.Status,
FlightDelaysValidation_Predicted2[,2])
```

```
##
```

```
## Call:
```

```
## roc.default(response = FlightDelaysValidation$Flight.Status,      predictor
= FlightDelaysValidation_Predicted2[, 2])
```


```
##
```

```
## Data: FlightDelaysValidation_Predicted2[, 2] in 190 controls
(FlightDelaysValidation$Flight.Status delayed) < 693 cases
(FlightDelaysValidation$Flight.Status ontime).
```

```
## Area under the curve: 0.6588
```

```
# Plot the ROC curve
```

```
plot.roc(FlightDelaysValidation$Flight.Status,  
FlightDelaysValidation_Predicted2[,2])
```



From the AUC value of 0.6588, it can be stated that the model is better than chance guessing; however, the model still has some room for improvement to increase this AUC value closer and closer to nearly 1.0.